

# CS 33 – Computer Organization

Week 0 Discussion  
Friday, 29 September 2017

Slides adapted from Uen-Tao Wang

# Contact Info

- TA: Brian Hill
- Email: [blhill@ucla.edu](mailto:blhill@ucla.edu)
- Office Hours: Tues/Wed 10:30-11:30am (tentative)
- TA Office Hours Location: BH 2432

# Schedule

- Administrative information
- Linux overview and accessing the SEASnet Linux servers
- C (aka unlearning C++)
- Binary representation
- Binary operators

# Course Administration

- Course Website: <http://web.cs.ucla.edu/classes/fall17/cs33/>
- Syllabus link available on course website
- Professor Eggert's Office Hours: Mon 2-3pm, Thur 10-11am
- Textbook: Computer Systems: A Programmer's Perspective (3rd edit), Randal Bryant & David O'Hallaron
  - Note: 2<sup>nd</sup> edition homework problems are different numbers than 3<sup>rd</sup>
- Grading:

• 5 Homeworks	5%	(1% each)
• 4 Labs	40%	
• 2 Midterms	25%	(12.5% each)
• 1 Final Exam	30%	

# Getting Started

- This class is based around C, not C++.
- As a result, you are highly recommended to ditch Visual Studio and work in a Linux environment, specifically the SEASnet Linux servers.
- Your assignments will be tested on the SEASnet Linux servers
  - `lnxsrv06`, `lnxsrv07`, `lnxsrv09` have newer version of gcc you'll want to use
- The class lectures are likely to be Linux heavy.
- Linux is love. Linux is life.



# Linux

# Getting Started: Accessing the SEASnet

- To login to SEASnet you need to be connected to wireless network on campus

OR

- Login with VPN Software

<https://www.it.ucla.edu/bol/services/virtual-private-network-vpn-clients>

# Getting Started: Accessing the SEASnet

- SSH stands for Secure Shell and is a protocol that is used to initiate text based access to a remote server.
- For Windows users:
  - PuTTY (<http://www.chiark.greenend.org.uk/~sgtatham/putty/download.html>): An SSH client
- For Mac/Linux:
  - SSH is a command that can be issued directly. Open a terminal (for Mac: Applications -> Utilities -> Terminal)

# Getting Started: Accessing the SEASnet

- For Windows users:
  - In “host name”, enter [username]@lnxsrv.seas.ucla.edu
- For Mac and Linux:
  - ssh [username]@lnxsrv.seas.ucla.edu
- If you are in the Henry Samueli School of Engineering, you should already have a SEASnet account. Otherwise go to the SEASnet office at 2684 Boelter Hall to get one.

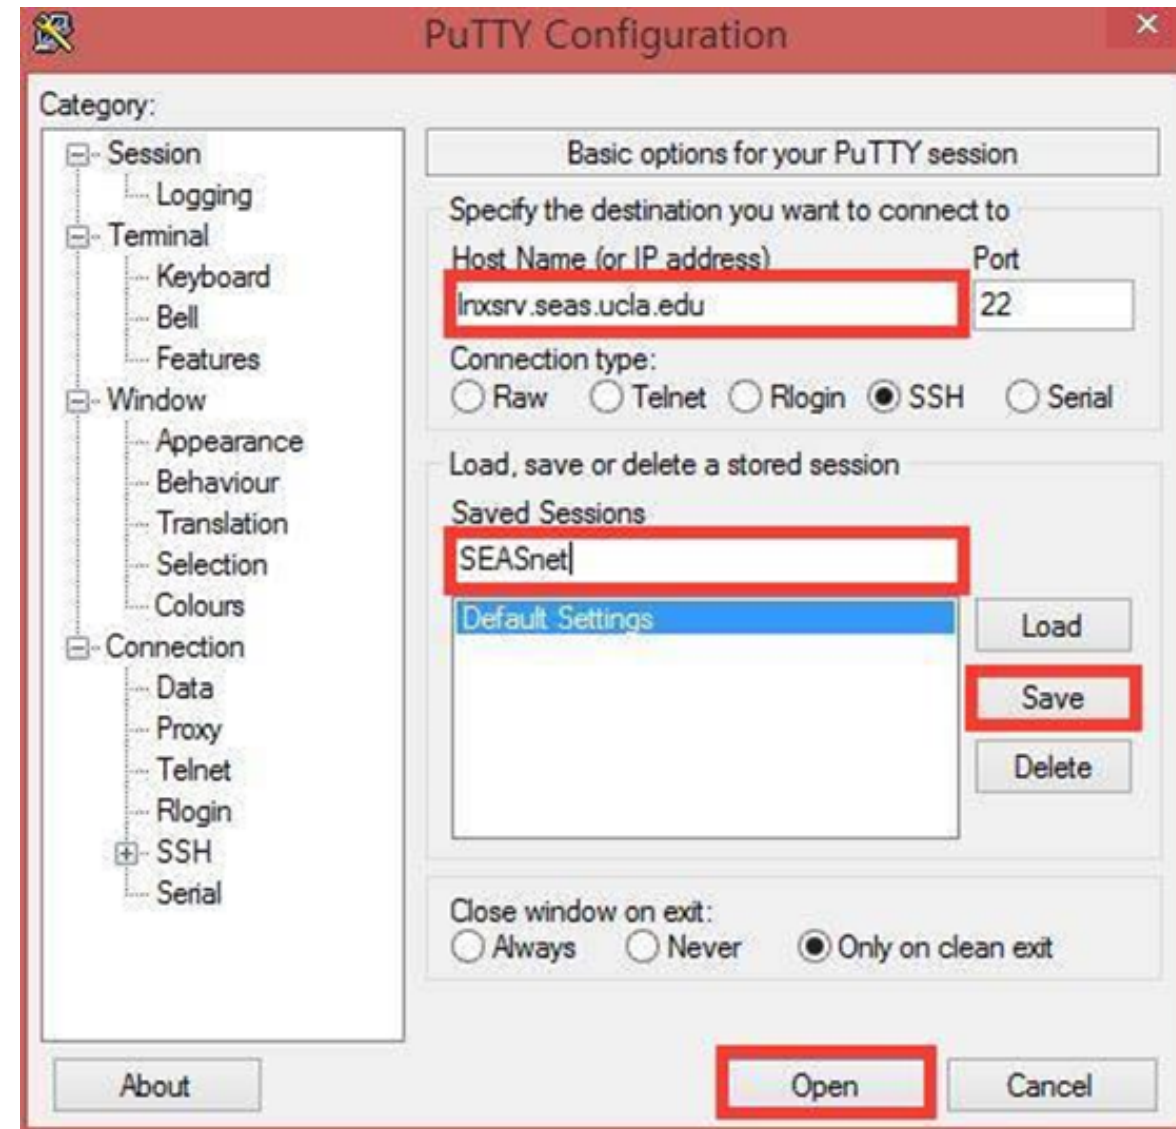


# PuTTY

## First Run

- Type *Inxsrv@seas.ucla.edu* for Host Name
- Type SEASnet for Saved Sessions
- Click Save
- Click Open
- Type your SEASnet username and password

Double-click SEASnet under Saved Sessions  
in the future



# Getting Started: Useful Linux Commands

- Typical Linux command format: `[command name] -X -Y -Z [argument]`
  - (X, Y, and Z are optional flags)
- Flags modify/specify the behavior of the command.
- Ex:
  - `ls` Lists files in current directory
  - `ls -l` Lists files in current directory, in “long” format
- **VERY** useful command: `man <command name>`
  - Opens manual page for the command
  - Ex: `man ls`

# Getting Started: Useful Linux Commands

Command	Example	Explanation
<code>pwd</code>	<code>pwd</code>	Print the working directory (current directory)
<code>ls [flags] [path to directory]</code>	<code>ls -l /tmp</code>	List contents of directory
<code>cd &lt;path to directory&gt;</code>	<code>cd /tmp</code>	Change working directory to argument
<code>mkdir [new directory name]</code>	<code>mkdir /tmp/test_folder</code>	Create new directory
<code>rm [flags] [file/directory]</code>	<code>rm -rf /tmp/test_folder</code>	Remove file/directory (be careful!!)
<code>exit</code>	<code>exit</code>	Exit terminal

# Getting Started: Useful Linux Commands

- Editing files. If you're interested familiarizing yourselves with Linux (which will have to happen eventually), it is recommended that you use “vim” or “emacs”.
  - vim text.txt
  - emacs text.txt
- Useful vim commands:

i	Insert mode (for editing text)
esc (Escape)	Get out of “insert mode”
:q	Exit vim
:wq	Write changes and exit vim
:q!	Do NOT save changes and exit vim

# Getting Started: Useful Linux Commands

- The standard Linux C compiler is gcc
  - `gcc main.c` (compile the file main.c into an executable file with default name "a.out")
  - `gcc main.c -o main` (compile the file main.c into an executable file called "main")
  - `gcc main.c -O2` (compile the file with optimizations, level 2)
  - `gcc -S main.c` (dump assembly code)
  - `gcc -E main.c` (show code after pre-processing)
- Executing executables
  - `./main` (executes the executable file called "main")

# C (as opposed to C++)

- In a (very simplified) nutshell, C++ is an extension to C.
- The syntax of the language is nearly identical, but you will find that C lacks certain features, namely the “Object Oriented” paradigm.
- Some features are analogous, but have different names.

# C (as opposed to C++)

- In C++:

- `for(int i = 0; i < size; i++)`

- By default, gcc uses a 1990's C standard which prohibits declarations in “for” loops. As a result, you will have to do either

- `int i;`

- `for(i = 0; i < size; i++)`

- Or explicitly use gcc to compile with a different C standard

- `gcc -std=c99 temp.c`

# C (as opposed to C++)

Dynamic memory allocation

In C++:

- `char * c_arr = new char[10];`
- `delete c_arr;`
- “new” allows you to specify repetitions of a specific data type.


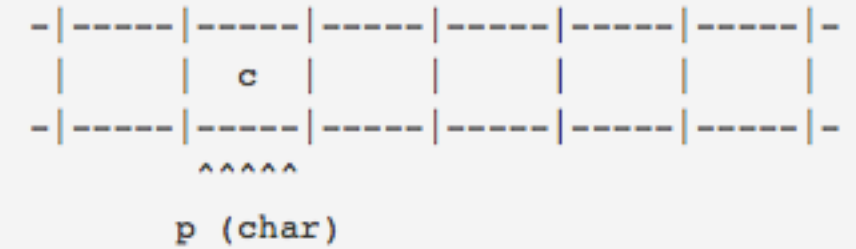


# C (as opposed to C++)

- In C, these declarations force you to be more specific. Instead of “new”, use “malloc” and instead of “delete”, use “free”.
  - `char * c_arr = (char *) malloc(sizeof(char) * 10);`
  - `free(c_arr);`
- Note: These are analogous but not the same.
- “malloc” and other “\_alloc” variations operate on the principle that you're specifying a specific amount of memory to allocate rather than a specific data type.

# C (as opposed to C++)

- Pointer casting: doesn't change address pointed to, but how we interpret data at that address
- Ex:
  - `char * c_arr = (char *) malloc(sizeof(char) * 10);`
- Casts void pointer (`void *`) to (`char *`)
- Example on right shows difference between casting pointer as (`char *`) vs (`int *`)



# C (as opposed to C++)

- Instead of:
  - `int x = 10;`
  - `cout << x;`
- You'll use “printf”
  - `printf(“hello”);`
  - `printf(“%d”, x);`
- `printf` takes in as the first parameter a string to print out that is populated with format codes that correspond to the remaining arguments.


Documentation: <http://www.cplusplus.com/reference/cstdio/printf/>

# -fwrapv and -ftrapv Flags

- `fwrapv` - instructs the compiler to assume that signed arithmetic overflow of addition, subtraction and multiplication wraps around using two's-complement representation. Enabled by default for Java.
- `ftrapv` - generates traps for signed overflow on addition, subtraction, multiplication operations

# Everything is bits

- Each bit is 0 or 1
- By encoding/interpreting sets of bits in various ways
  - Computers determine what to do (instructions)
  - ... and represent and manipulate numbers, sets, strings, etc...
- Why bits? Electronic Implementation
  - Easy to store with bistable elements
  - Reliably transmitted on noisy and inaccurate wires



# For example, can count in binary

- Base 2 Number Representation
  - Represent  $15213_{10}$  as  $11101101101101_2$
  - Represent  $1.20_{10}$  as  $1.0011001100110011[0011]....._2$
  - Represent  $1.5213 \times 10^4$  as  $1.1101101101101_2 \times 2^{13}$

# Encoding Integers

Unsigned

$$B2U(X) = \sum_{i=0}^{w-1} x_i \cdot 2^i$$

```
short int x = 15213;  
short int y = -15213;
```

Two's Complement

$$B2T(X) = -x_{w-1} \cdot 2^{w-1} + \sum_{i=0}^{w-2} x_i \cdot 2^i$$

Sign  
Bit



- C short 2 bytes long

	Decimal	Hex	Binary
<b>x</b>	15213	3B 6D	00111011 01101101
<b>y</b>	-15213	C4 93	11000100 10010011

- Sign Bit
  - For 2's complement, most significant bit indicates sign
    - 0 for nonnegative
    - 1 for negative

# Two's complement Encoding Example (Cont.)

**x =**           15213: 00111011 01101101  
**y =**           -15213: 11000100 10010011

Weight	15213		-15213	
1	1	1	1	1
2	0	0	1	2
4	1	4	0	0
8	1	8	0	0
16	0	0	1	16
32	1	32	0	0
64	1	64	0	0
128	0	0	1	128
256	1	256	0	0
512	1	512	0	0
1024	0	0	1	1024
2048	1	2048	0	0
4096	1	4096	0	0
8192	1	8192	0	0
16384	0	0	1	16384
-32768	0	0	1	-32768
Sum	15213		-15213	



# Numeric Ranges

## Unsigned Values

- $UMin = 0$   
000...0
- $UMax = 2^w - 1$   
111...1

- **Two's Complement Values**

- $TMin = -2^{w-1}$ 
  - 100...0
- $TMax = 2^{w-1} - 1$ 
  - 011...1

- **Other Values**

- Minus 1
  - 111...1

## Values for $W = 16$

	Decimal	Hex	Binary
UMax	<b>65535</b>	FF FF	11111111 11111111
TMax	<b>32767</b>	7F FF	01111111 11111111
TMin	<b>-32768</b>	80 00	10000000 00000000
-1	<b>-1</b>	FF FF	11111111 11111111
0	<b>0</b>	00 00	00000000 00000000

# Values for Different Word Sizes

	W			
	8	16	32	64
UMax	255	65,535	4,294,967,295	18,446,744,073,709,551,615
TMax	127	32,767	2,147,483,647	9,223,372,036,854,775,807
TMin	-128	-32,768	-2,147,483,648	-9,223,372,036,854,775,808

## Observations

- $|TMin| = TMax + 1$
- $UMax = 2 * TMax + 1$

## C Programming

- `#include <limits.h>`
- Declares constants, e.g.,
  - `ULONG_MAX`
  - `LONG_MAX`
  - `LONG_MIN`
- Values platform specific

# Unsigned & Signed Numeric Values

$X$	$B2U(X)$	$B2T(X)$
0000	0	0
0001	1	1
0010	2	2
0011	3	3
0100	4	4
0101	5	5
0110	6	6
0111	7	7
1000	8	-8
1001	9	-7
1010	10	-6
1011	11	-5
1100	12	-4
1101	13	-3
1110	14	-2
1111	15	-1

## Equivalence

- Same encodings for nonnegative values

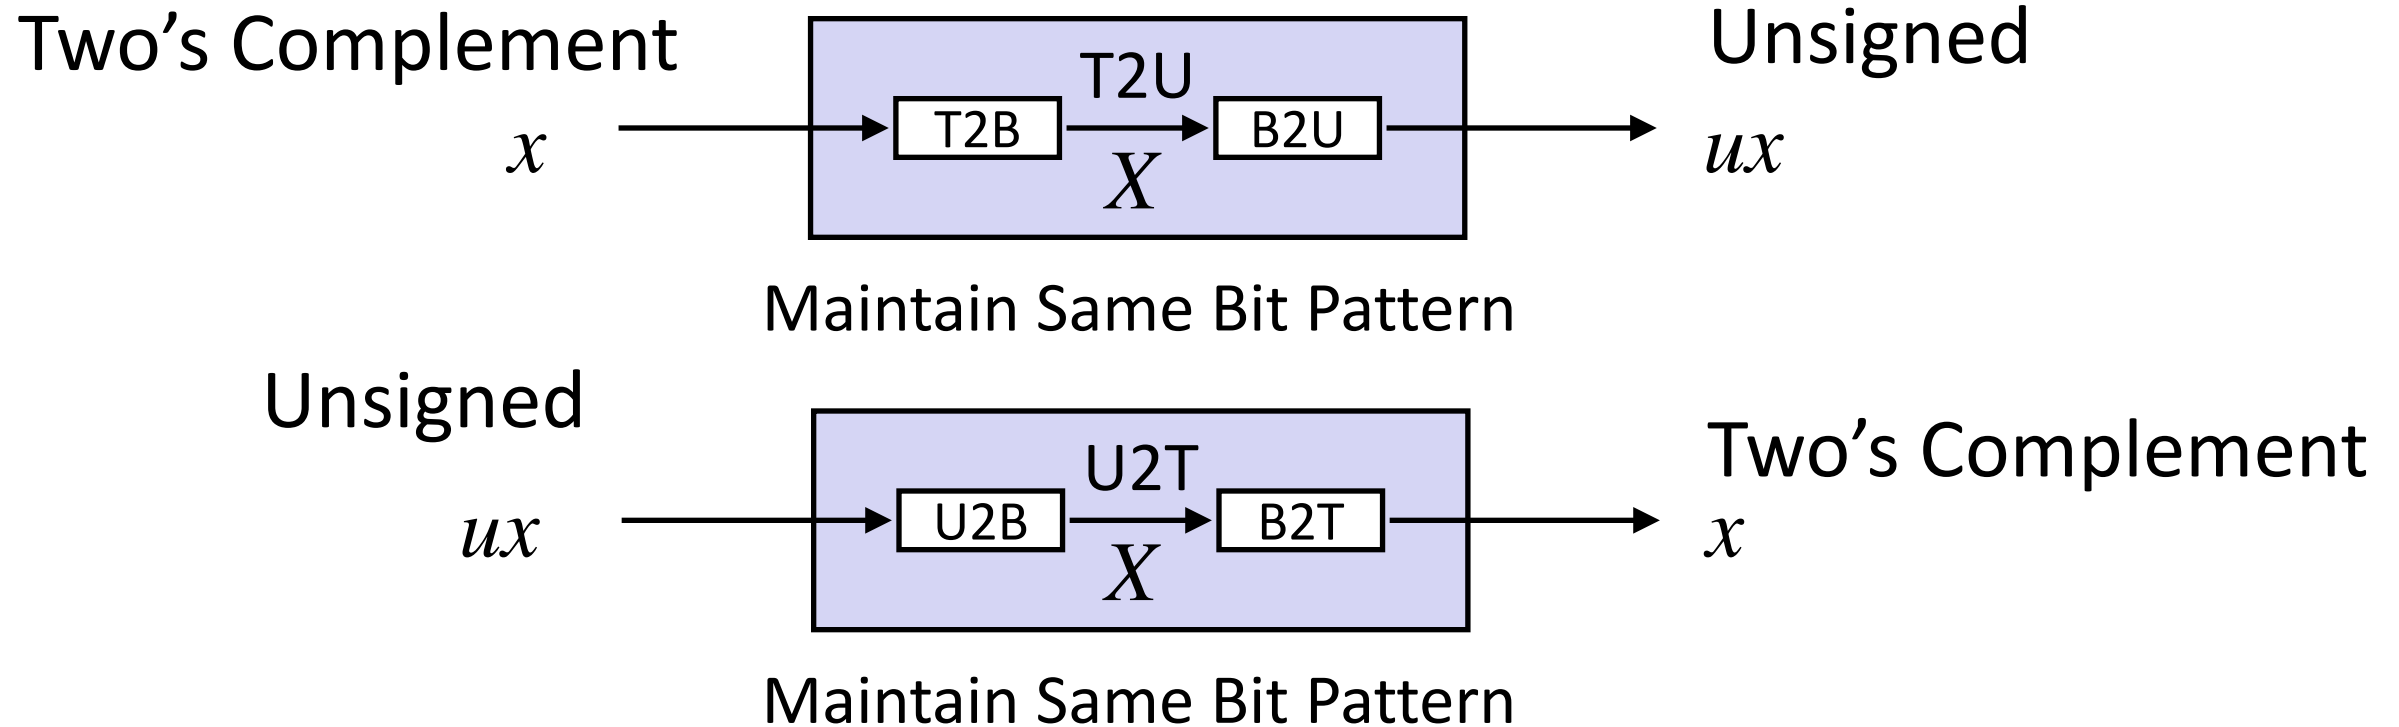
## Uniqueness

- Every bit pattern represents unique integer value
- Each representable integer has unique bit encoding

## ⇒ Can Invert Mappings

- $U2B(x) = B2U^{-1}(x)$ 
  - Bit pattern for unsigned integer
- $T2B(x) = B2T^{-1}(x)$ 
  - Bit pattern for two's comp integer

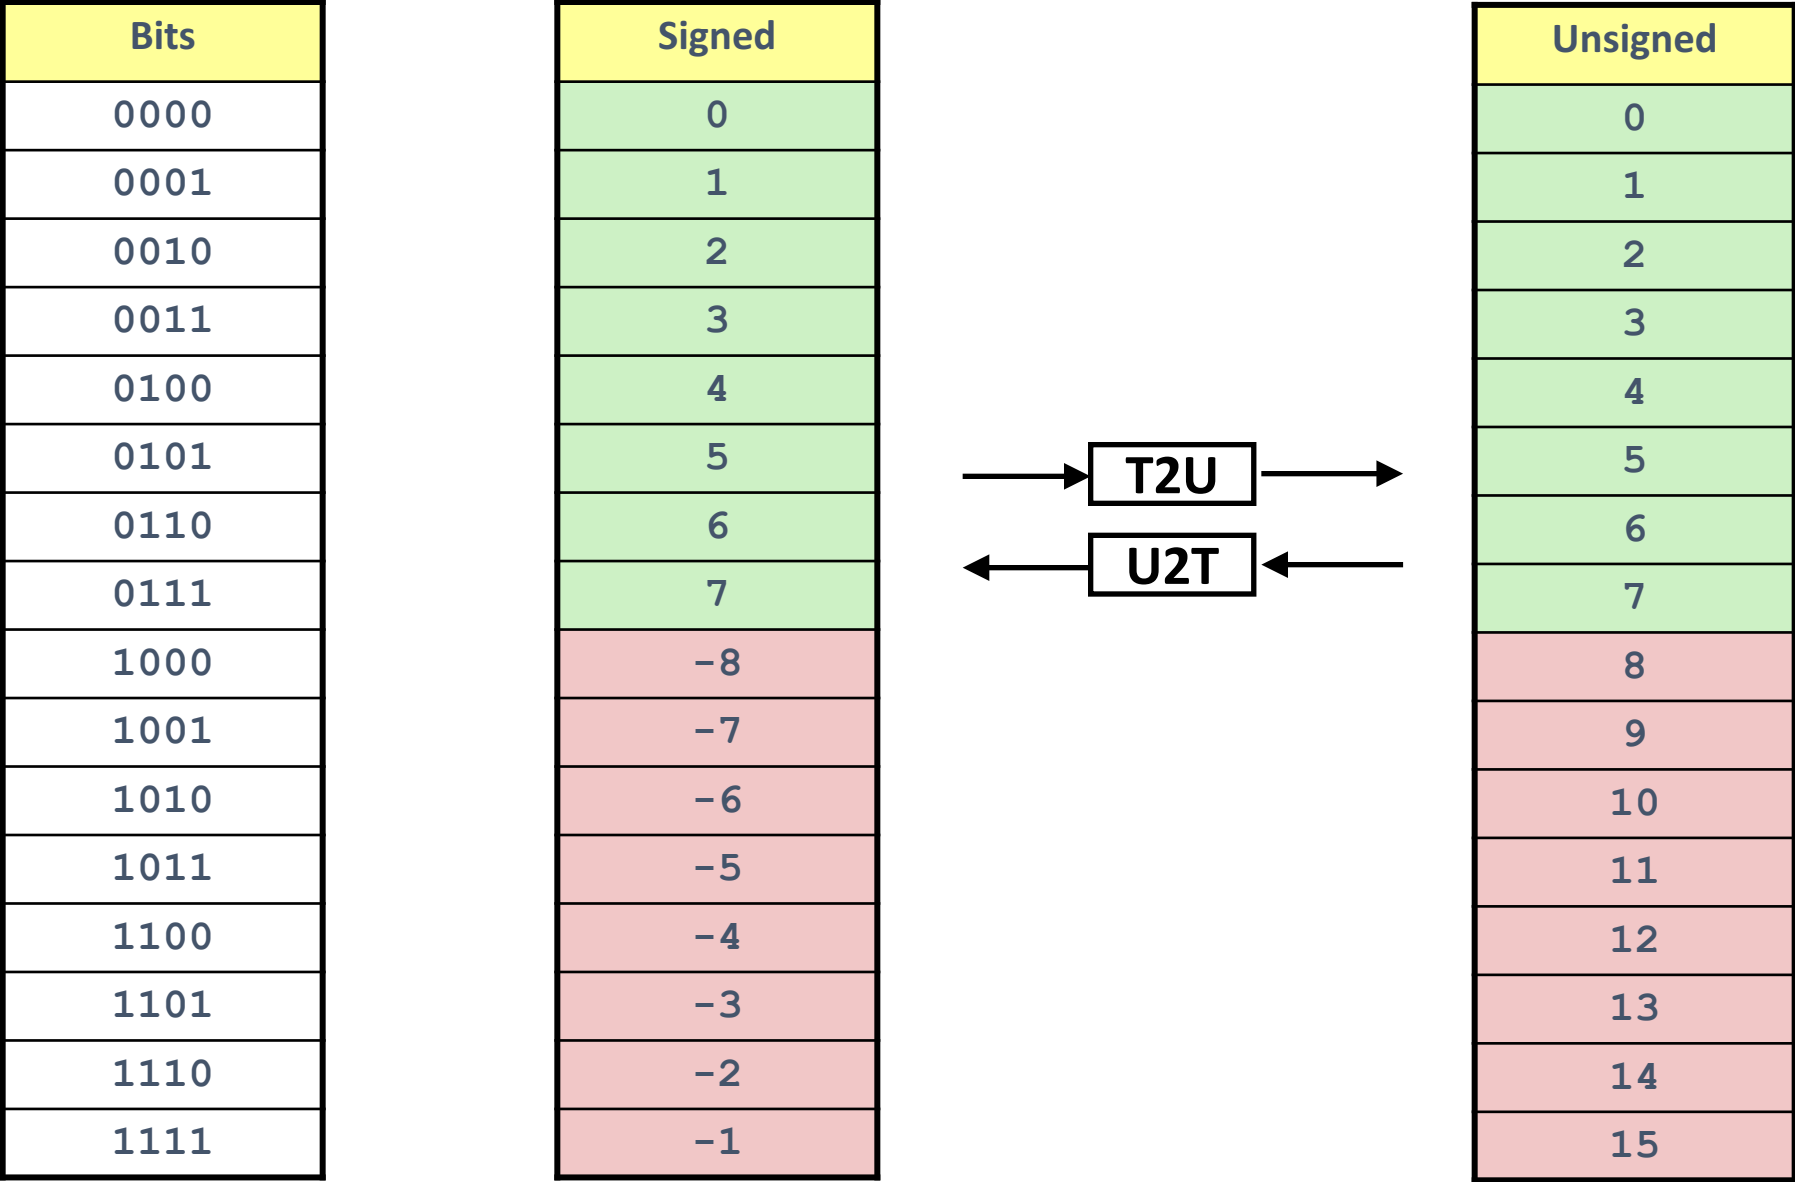
# Mapping Between Signed & Unsigned





Mappings between unsigned and two's complement numbers:

**Keep bit representations and reinterpret**


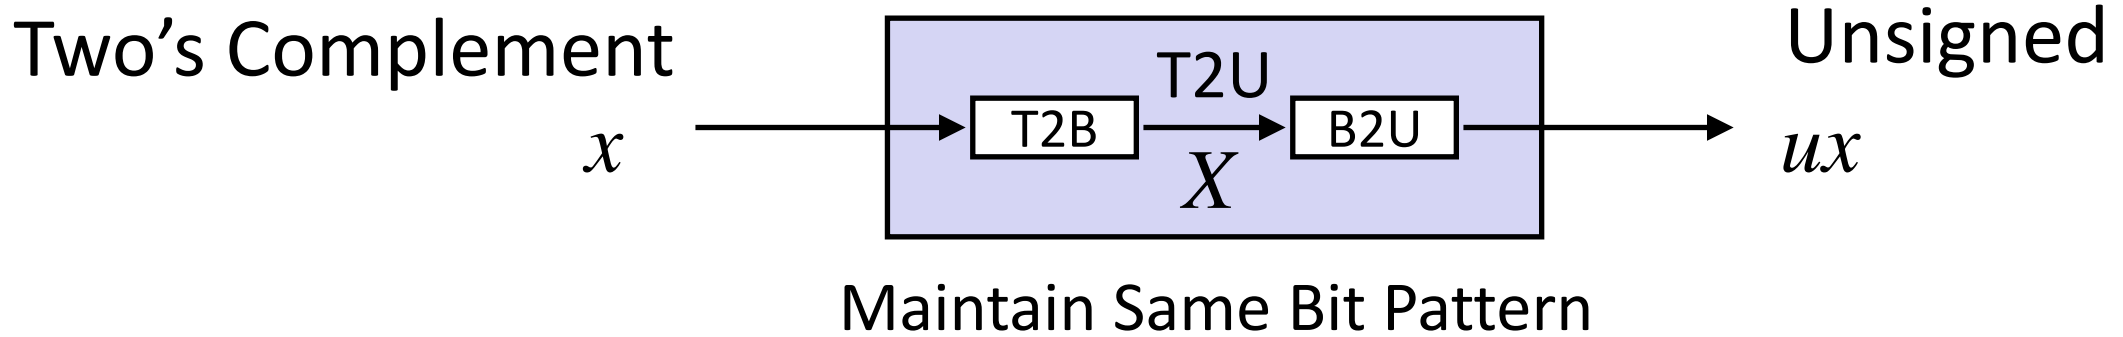
# Mapping Signed $\leftrightarrow$ Unsigned



# Mapping Signed $\leftrightarrow$ Unsigned

Bits	Signed		Unsigned
0000	0		0
0001	1		1
0010	2		2
0011	3		3
0100	4		4
0101	5		5
0110	6		6
0111	7		7
1000	-8		8
1001	-7		9
1010	-6		10
1011	-5		11
1100	-4		12
1101	-3		13
1110	-2		14
1111	-1		15


# Relation between Signed & Unsigned



# Conversion Visualized

## 2's Comp. → Unsigned

- Ordering Inversion
- Negative → Big Positive





# Signed vs. Unsigned in C

## Constants

- By default are considered to be signed integers
- Unsigned if have “U” as suffix  
`0U, 4294967259U`

## Casting

- Explicit casting between signed & unsigned same as U2T and T2U  
`int tx, ty;`  
`unsigned ux, uy;`  
`tx = (int) ux;`  
`uy = (unsigned) ty;`
- Implicit casting also occurs via assignments and procedure calls  
`tx = ux;`  
`uy = ty;`

# Casting Surprises

## Expression Evaluation

- If there is a mix of unsigned and signed in single expression,  
***signed values implicitly cast to unsigned***
- Including comparison operations  $<$ ,  $>$ ,  $==$ ,  $<=$ ,  $>=$
- Examples for  $W = 32$ : **TMIN = -2,147,483,648 , TMAX = 2,147,483,647**

# Casting Surprises

Constant <sub>1</sub>	Constant <sub>2</sub>	Relation	Evaluation
0	0U	==	unsigned
-1	0	<	signed
-1	0U	>	signed
2147483647	-2147483647-1	>	signed
2147483647U	2147483647U	<	unsigned
-1	-2	>	signed
(unsigned)-1	-2	>	unsigned
2147483647	2147483648U	<	unsigned
2147483647	(int) 2147483648U	>	signed


# Sign Extension

## Task:

- Given  $w$ -bit signed integer  $x$
- Convert it to  $w+k$ -bit integer with same value

## Rule:

- Make  $k$  copies of sign bit:
- $X' = x_{w-1}, \dots, x_{w-1}, x_{w-1}, x_{w-2}, \dots, x_0$



# Sign Extension Example

```
short int x = 15213;  
int      ix = (int) x;  
short int y = -15213;  
int      iy = (int) y;
```

	Decimal	Hex	Binary
<b>x</b>	15213	3B 6D	00111011 01101101
<b>ix</b>	15213	00 00 3B 6D	00000000 00000000 00111011 01101101
<b>y</b>	-15213	C4 93	11000100 10010011
<b>iy</b>	-15213	FF FF C4 93	11111111 11111111 11000100 10010011

Converting from smaller to larger integer data type

C automatically performs sign extension

# Summary: Expanding, Truncating: Basic Rules

## **Expanding (e.g., short int to int)**

- Unsigned: zeros added
- Signed: sign extension
- Both yield expected result

## **Truncating (e.g., unsigned to unsigned short)**

- Unsigned/signed: bits are truncated
- Result reinterpreted
- Unsigned: mod operation
- Signed: similar to mod
- For small numbers yields expected behavior

# Signed Binary - Two's Complement

How do we represent negative numbers?

The two's complement of a number is technically its value subtracted from  $2^N$ .

In two's complement, most bits have the same contribution as in unsigned.

The value of the  $i$ -th bit is  $2^i$  (assuming  $i$  starts from 0).

However, the most significant bit of an  $N$  bit number has a value of  $-2^{(N-1)}$  instead of  $2^{(N-1)}$

# Signed Binary: Two's Complement

- Assume we're dealing with four bit numbers.
- Consider the unsigned binary number 1010:
  - $1010 = 1*2^3 + 0*2^2 + 1*2^1 + 0*2^0 = 10$
- Now consider the signed binary number 1010:
  - $1010 = 1*(-(2^3)) + 0*2^2 + 1*2^1 + 0*2^0 = -6$
- Same sequence of bits, but depends on how we interpret!



# How to convert between negative and positive

- The method: take the bitwise inverse and add one. Consider 0101 (5).

- 1. 0101
- 2. Bitwise inverse of 0101 = 1010.
- 3.  $1010 + 0001 = 1011$ .
- 4. Confirm:  $1011 = -(2^3) + 2^1 + 2^0 = -5$

# Signed Binary: Notes

- The value of a signed binary number depends on the number of bits there are.
  - Four bit signed: 1111 = -1
  - Five bit signed: 01111 = 15
- An N-bit signed binary number has  $2^N$  possible values with a range of  $[-2^{(N-1)}, 2^{(N-1)} - 1]$
- REMEMBER THIS: The range of a two's complement signed binary number is **not** symmetrical around 0.
- Henceforth all signed binary is two's complement unless otherwise specified.

# Binary arithmetic

- What does it mean to add bits?
- The idea is the same as in decimal. Let's try with unsigned.

	0001		0001		0001		0011
+	0010		+ 0001		+ 0111		+ 0111
	-----		-----		-----		-----
	0011		0010		1000		1010

Note that adding two or three 1 bits will produce a carry bit that must be added to the next bit over.

# Binary arithmetic

- How about subtraction? You can generalize the decimal method for binary, but...
- The simplest way to do  $X - Y$  is to do  $X + (-Y)$ .
- Take  $0110$  (6)  $- 0010$  (2)
- This becomes  $0110$  (6)  $+ 1110$  (-2)

$$\begin{array}{r} 0110 \\ + 1110 \\ \hline 0100 \end{array}$$

# Unsigned overflow in C

- With signed arithmetic, we saw that a carry out bit was completely valid, but what about unsigned?

- Say we have 4-bit numbers and we add  $6 + 12$ .

- $6 = 0110$ ,  $12 = 1100$

$$\begin{array}{r} 0110 \\ + 1100 \\ \hline 10010 = 18 \end{array}$$

- ...but this requires 5 bits to represent. We only have 4.
- How does the wise and venerable C respond?

# Unsigned overflow in C

- Just drop bits.
- If an unsigned operation of an n-bit number requires more than n bits, the resulting number will consist only of the n least significant bits.
  - Ex. In the previous example:
  - $0110 + 1100 = (1) 0010$ , the leading one is dropped and instead of the right answer of 18, you get the incredibly wrong answer of 2
- More formally, if you have n-bits, the computation of  $x + y$  is  $(x + y) \% 2^n$ .
  - Ex.  $(6 + 12) \% 2^4 = 18 \% 16 = 2$

# Unsigned overflow in C

- This is also true of unsigned multiplication overflow:
- If we have 4 bits,  $6 * 12 = 72$
- In binary, 72 is 1001000.
- Truncate bits beyond 4 and the result is 1000 = 8.
- $72 \% 2^4 = 8$
- Keep in mind, this is for unsigned numbers only.
- Let's not think about signed numbers for now.

# Datatypes in C

- Each native datatype in C is expressed by a sequence of bits.
- Simple data types such as ints and shorts come in unsigned and signed variants where signed is the default (ie. int is actually a signed int)
- However, the number of bits used to express these numbers differs depending on whether the processor is 32 or 64-bit.
- ...but more on the processor definitions later.



# Datatypes in C

- char/unsigned char : 8-bits
- short/unsigned short : 16-bits
- int/unsigned (int) : 32-bits

Here's where it gets weird.

- In 32-bit machines:
  - long/unsigned long : 32-bits
  - long long/unsigned long long (proof that a five year old named these) : 64-bits
- In 64-bit machines:
  - long/unsigned long : 64-bits
  - long long/unsigned long long (ugh) : 64-bits

# Boolean Operators

- Boolean operators operate on a single bit.
- AND : &
  - Result is 1 if both inputs are 1.
- OR : |
  - Result is 1 if either of the inputs are 1.
- XOR : ^
  - Result is 1 if one input is 1 and the other is 0
- NOT : ~
  - Result is 1 if the input is 0.

# Bitwise Operators

- Bitwise operators perform repeated boolean operations on each bit of a number or pair of numbers.
- Bitwise invert (not the same as logical invert or '!')
  - $\sim(1011) = 0100$
- Bitwise AND/OR (not the same as logical AND/OR or `&&/||`)
  - $1010 \& 1100 = 1000$
  - $1010 | 1100 = 1110$
- Bitwise XOR
  - $1010 \wedge 1100 = 0110$

# Bitwise Operators

- Left shift/right shift (arithmetic vs logical)
- Left shift
  - $0111 \ll 1 = 1110$
- Right shift
  - $1011 \gg 1 = 0101$  (logical)
  - $1011 \gg 1 = 1101$  (arithmetic)
- Why two different right shifts?

# Logical Operators

- Where bitwise operators operate on each individual bit of a number, logical operators operate on the number as a whole
  - `||`, `&&`, `!`
- To invert a bit sequence `x`, you would use `~x`.

What happens if you use the logical invert `!`?

- `!(1010) = 0`
- `!(0111) = 0`
- `!(0) = 1`

# Logical Operators

- What happens when you use logical operators on numbers?  
1011 && 1100?
- Non-zero numbers are interpreted as 1 and 0 is interpreted as 0.  
1011 && 1100 <=> 1  
1011 && 0 <=> 0  
1011 || 0 <=> 1

# Useful Tips

x is a bit vector  
if(x == 0)  
 return 0;  
else  
 return 1;

OR

return !!x;



a, b, and c are bits

```
if(a)
    return b;
else
    return c;
```

OR

```
return (a & b) | (~a & c)
```

# De-Morgan's Law

$$a \& b = \sim(\sim a \mid \sim b)$$

$$a \mid b = \sim(\sim a \& \sim b)$$

# Multiplication by Shifting

- Consider the 4-bit unsigned number 0110.
  - $0110 = 2^3 * 0 + 2^2 * 1 + 2^1 * 1 + 2^0 * 0 = 2^2 + 2^1 = 6$
- $0110 \ll 1 = 1100$ 
  - $1100 = 2^3 * 1 + 2^2 * 1 + 2^1 * 0 + 2^0 * 0 = 2^3 + 2^2$
  - $2^3 + 2^2 = 2 * (2^2 + 2^1) = 12$
- $x \ll n = x * 2^n$

# Multiplication by Shifting

How can we think of multiplying two arbitrary (ie non-powers of two) numbers in binary?

$$\begin{aligned}0110 * 1011 & (= 6 * 11 = 66) \\&= 0110 * (1000 + 0010 + 0001) \\&= 0110 * 1000 + 0110 * 0010 + 0110 * 0001 \\&= 0110 \ll 3 + 0110 \ll 1 + 0110 \\&= 0110000 + 01100 + 0110 \\&= 1000010\end{aligned}$$

# Division by shifting

By the same logic, this ought to work for division right?

Consider 4-bit unsigned:

- $1100 = 12$
- $1100 \gg 1 = 0110 = 6$

Consider 4-bit signed:

- $1100 = -4$
- $1100 \gg 1 = 0110 = 6$  (??)

# Division by shifting

- Previously, we tried logical right shifting (shift in zeros, but that didn't seem to pan out). This is where arithmetic right shifting steps in.
- Consider 4-bit signed:
  - $1100 = -4$
  - $1100 \gg 1 = 1110 = -2$
- Logical shifting maintains correct values for unsigned operations while arithmetic shifting maintains correct values for signed operations.

# Division by shifting

- Consider the 4-bit signed number 1101.
  - $1101 = -2^3 * 1 + 2^2 * 1 + 2^1 * 0 + 2^0 * 1 = -(2^3) + 2^2 + 2^0 = -3$
- $1101 \gg 1 = 1110$ 
  - $1110 = -2^3 * 1 + 2^2 * 1 + 2^1 * 1 + 2^0 * 0 = -2$
- $-3 / (\text{integer}) 2 = -1$ , not  $-2$
- How do you resolve this?

# Access specific bits

Say you have the binary value 1010 and you only want to consider bits 1 and 2, that is, you want to transform 1010 into 0010.

```
    1010
&    0110
-----
    0010
```



# References

- Slides modified from DJ Kim, UT Wang and Shikhar Malhotra
- <http://www.cs.cmu.edu/afs/cs/academic/class/15213-f15/www/schedule.html>

Thank You